Module 1

IMPORTANT

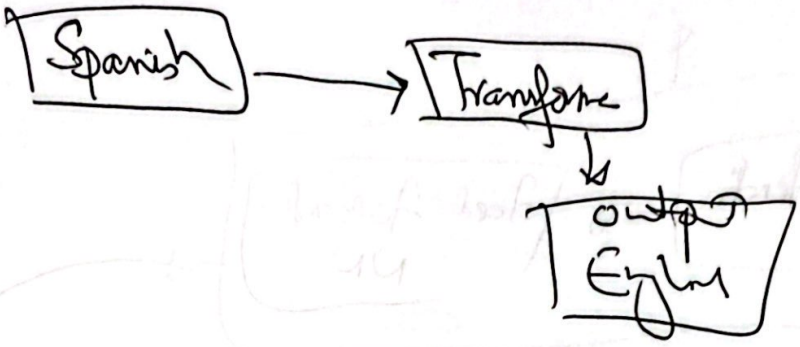

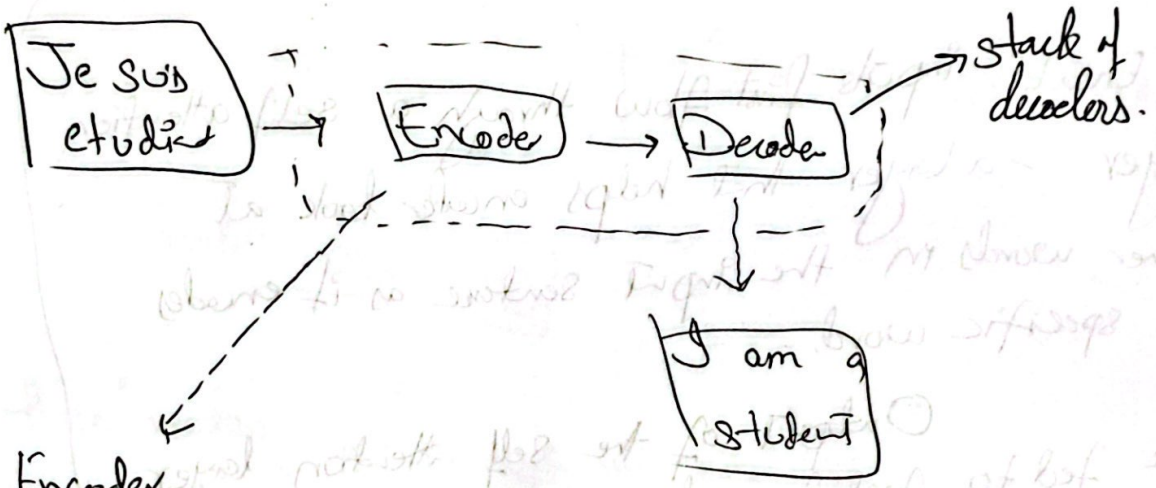

Encoder
↳ is actually a stack of encoders, no. 6
↳ hyperparameter
↳ encoder: 6
↳ decoder: 6

output of final encoder will be input to all decoders.

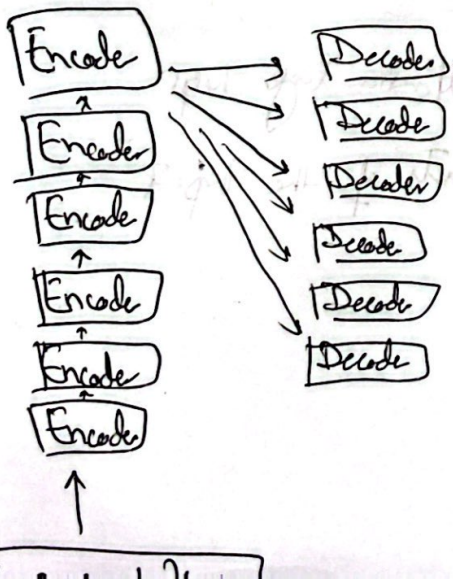

Inside Each encoder

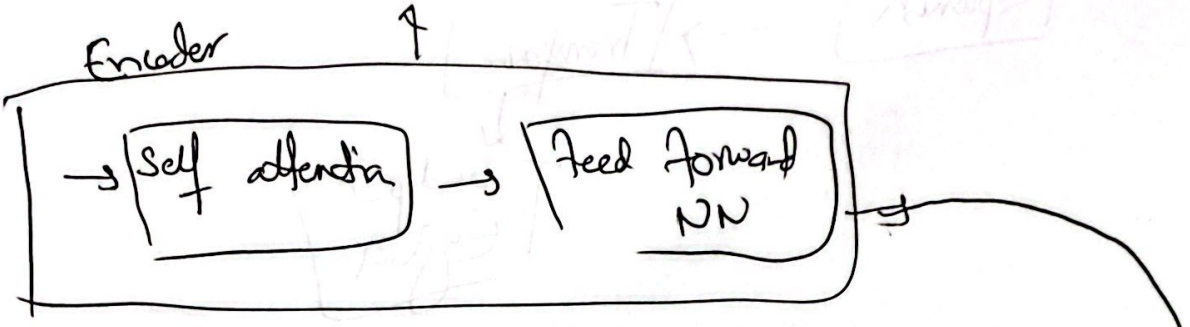

Encoder's inputs first flow through a self attention layer - a layer that helps encoder look at other words in the input sentence as it encodes a specific word.

Output of the self-attention layer are fed to feed forward NN. The exact same feed-forward ~~Net~~ NN is applied to each position.

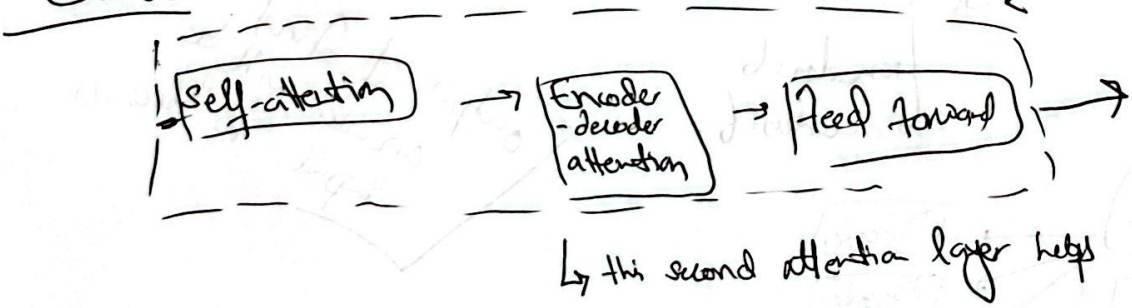

↳ this second attention layer helps the decoder focus on relevant parts of the input sentence.

Attention

Attention v/s Self Attention.

input: Je Suis etudiant

turn each word into vector.

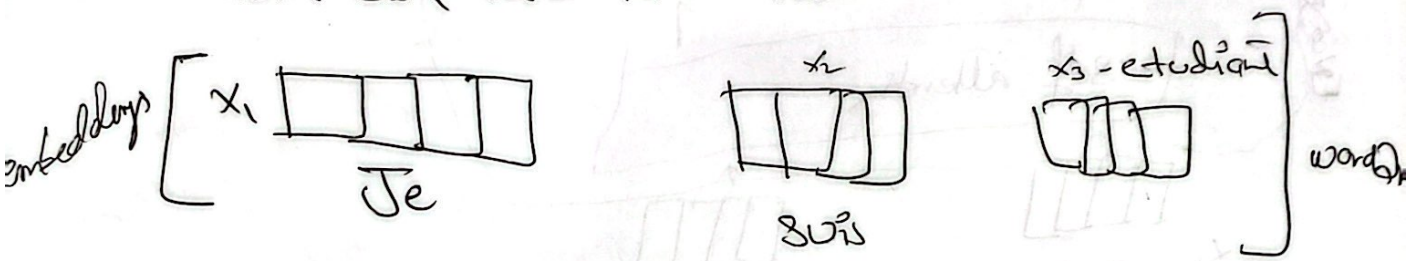

vector size - 512

these vectors are passed to self attention layer

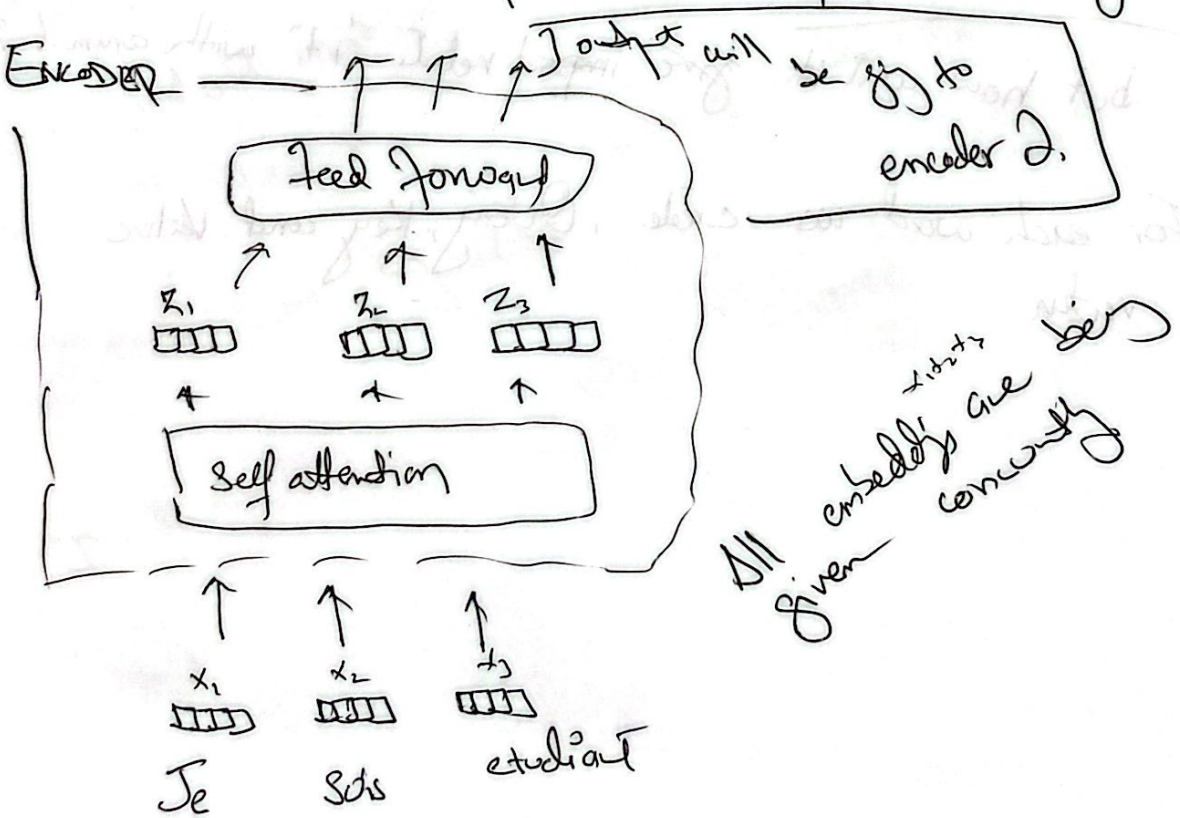

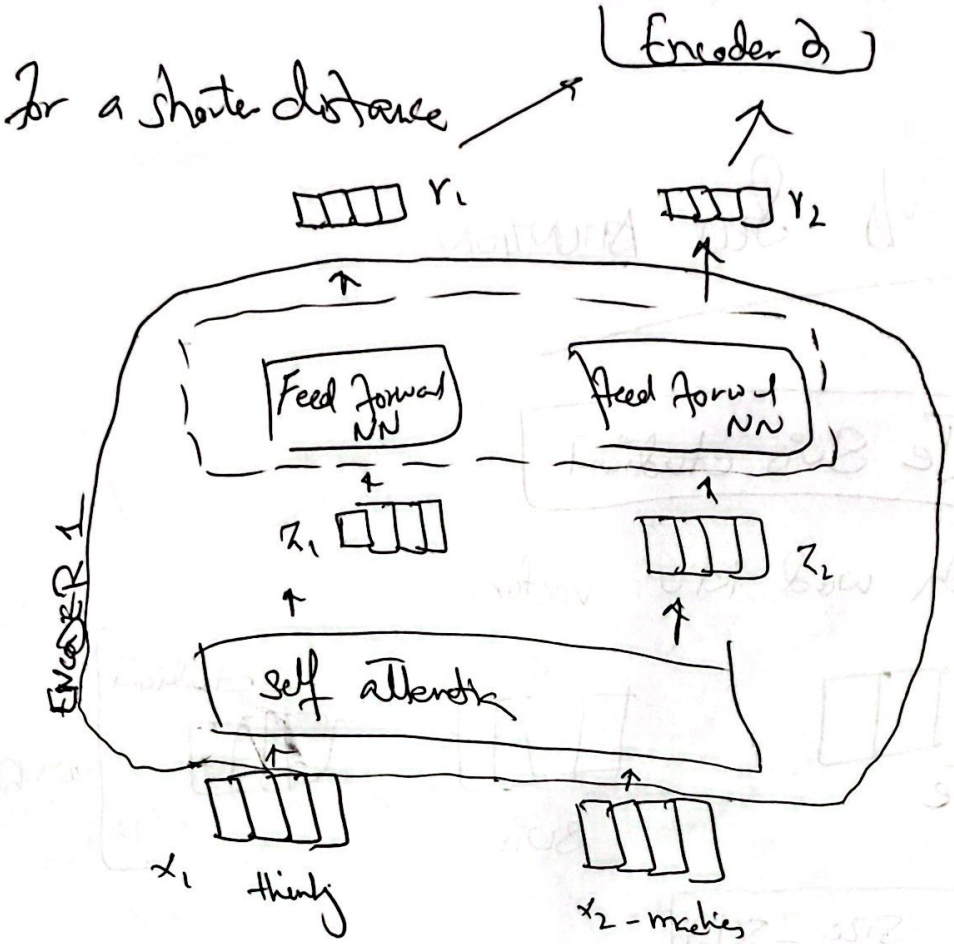

Self-attention,

but how will it ~~give~~ ~~input~~ relate 'it' with animal?

For each word we create, Query, Key and Value vector

Attention

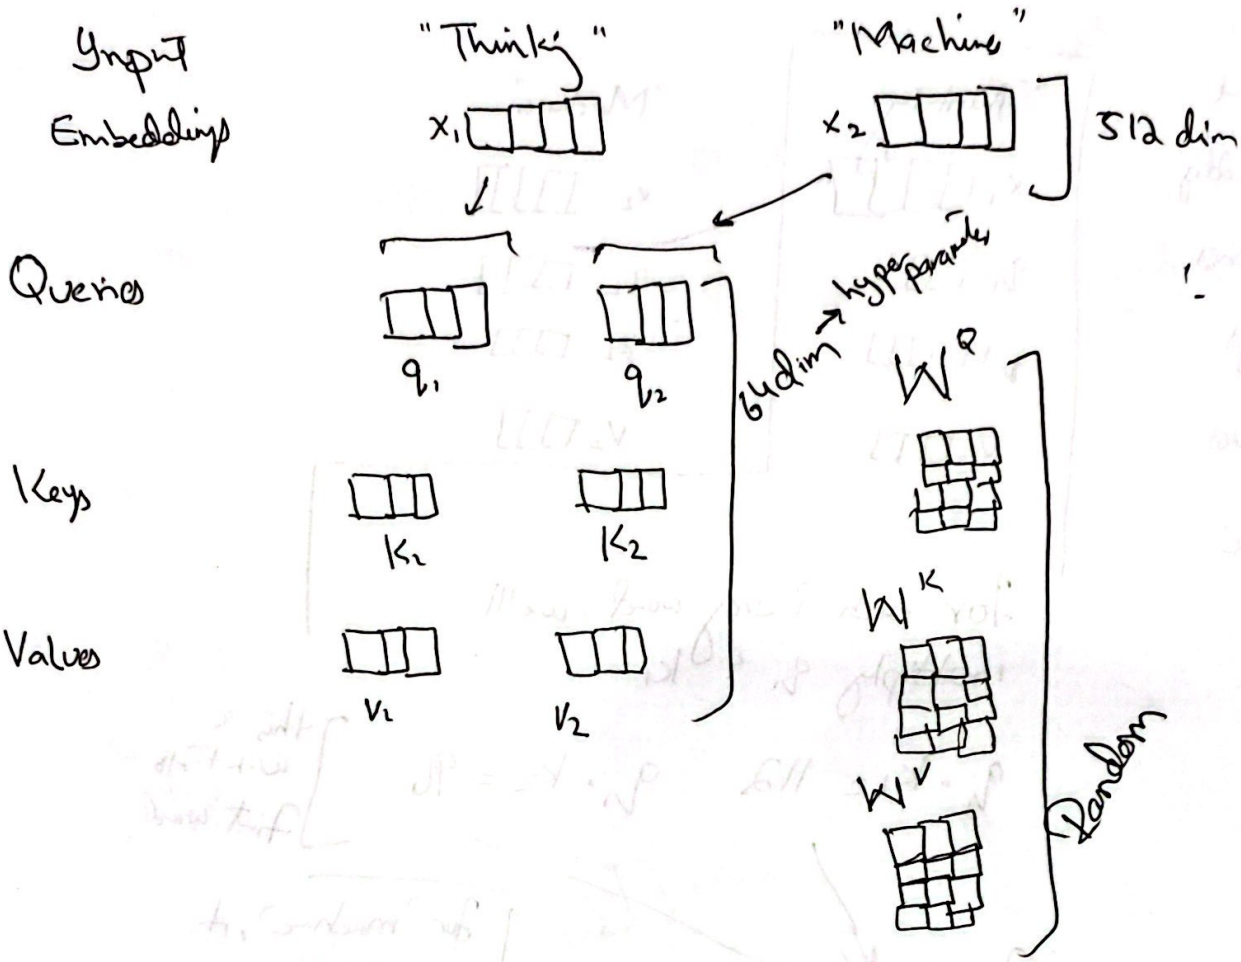

W^Q, W^K, W^V are 3 weights we randomly initialized. It will be updated based on back propagation.

So, first

$$q_1 = x_1 \cdot W^Q \qquad q_2 = x_2 \cdot W^Q$$] 64-d

why multiplied? input multiplied by weight + bias $Wx+b$ this how activate this how ANN operate.

Same for W^K & W^V

$$k_1 = x_1 \cdot W^K$$

$$v_1 = x_1 \cdot W^V$$

So,

Input Embedding
Queries
Keys
Values
Score

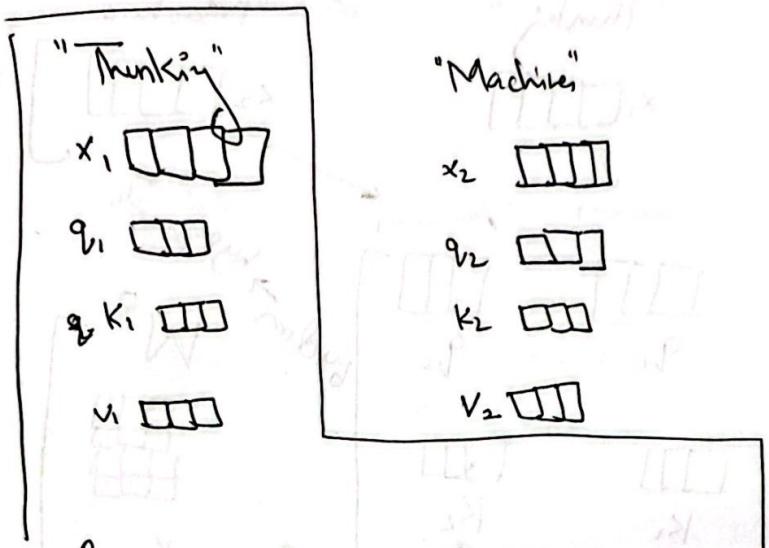

for each & every word, we'll multiply q_1 & k_1

$q_1 \cdot k_1 = 112$ $q_1 \cdot k_2 = 96$] this is w.r.t to first word.

Scores

for 'machine', it would be
$q_2 \cdot k_1 =$
$q_2 \cdot k_2 =$

We need to find out most related words.

Divide score by $\sqrt{64}$ — dim of q_1, keys : d_k

Divide by $\sqrt{64}$ ($\sqrt{d_k}$) = 112/8 = 14 ~~0.88~~ 96/8 = 12 ~~0.12~~

Softmax = 0.88 "thinking" 0.12 "machines" to set prob

Attention

Softmax 0.88 0.12

Softmax × value → multiply by V_1, V_2

V_1 × softmax

V_2 × softmax → to keep select words only → to give low prob to less relevant words

Sum : z_1 + $V_1 + V_2$

→ output of "this".

Matrix multiplication

$X \times W^Q = Q$

X: 2×4 (4 inputs, 2 words)

W^Q: 4×3

Q: 2×3

$X \times W^K = K$

$X \times W^V = V$

No window size because all words are being passed parallely.

Combining

$$\text{softmax}\left(\frac{Q \times K^{T}}{\sqrt{d_k}}\right) V = Z$$

(transposed)

Problem?

↳ We are using one W^Q, one W^K, one W^V for all words. This is single head attention.

Same W^K, W^Q, W^V for all words

but problem lies when there might might be important context.

So, we might use multiple weights

MULTI-HEAD Attention

↳ to find importance of other words too

- mean not only finding single word but multiple relevant words

Attention head 0 [Q_0 W_0^Q, K_0 W_0^K, V_0 W_0^V]

head 1 [Q_1 W_1^Q, K_1 W_1^K, V_1 W_1^V]

Different Heads for Different words

Attn 5

eight heads
Each eight diff weight matrices.
will give 8 z matrices.

So now

X
"thirty macha

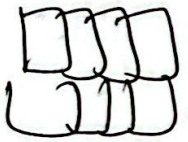

→

Attn head 0 → z_0 (3 × 2)
head 1 → z_1
head 2 → z_2
3 ⇒ z_3
4 → z_4
5 → z_5
6 → z_6
7 → z_7

But we are supposed to get 1 z value not 8.

1- we'll concatenate all Zs

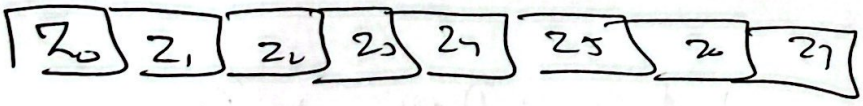

2- inside another W^0 matrix → would be updated in back propgn
W^0 will be multiplied with concated Z

result would be 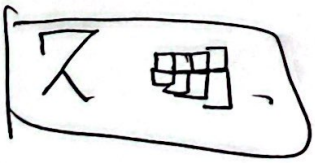

→ POSITIONAL ENCODING ~~before~~ after converting words into embeddings to preserve order.

INPUT	Je	SUIS	ETUDIANT
Embeddings	x_1	x_2	x_3
+	+	+	+
Positional encoding	t_1	t_2	t_3
Embeddings with time signal	x_1	x_2	x_3

t_1 would be nearer to t_2 than t_3.

This happens before passing to encoder

After this

In encoder, befor passing to feed forward ~~Each encoder~~ and after self attention, we'll be ~~perf~~ applying normalization and a residual connection.

Like in RNN, you are skipping one step.

If i.e if self attention is not useful, you are skipping it that's residual network

In add & Normalize

$$\text{layer norm}\left(X + Z\right)$$

X: input, Z: encoder output (Z)

In Decoder

Encoder Decoder Attention

- Save only output of encoder
- passed into the layer of decoder

input - all words at once

output - one word at a time.

- after each word predicts that word output will be give input to decoder. It won't stop until <EOS>.